Linux Kernel GDB tracepoint module (KGTP)

Update in 2014-01-01

Table of contents

3

About this document

https://code.google.com/p/kgtp/wiki/HOWTO is the the last version of this document in HTML format.

https://raw.github.com/teawater/kgtp/master/kgtp.pdf is the the last version of this document in PDF format.

https://raw.github.com/teawater/kgtp/release/kgtp.pdf is the the last release version of this document in PDF format.

What is KGTP

KGTP is a comprehensive dynamic tracer for analysing Linux kernel and application (including Android) problems on production systems in real time.

To use it, you don't need patch or rebuild the Linux kernel. Just build KGTP module and insmod it is OK.

It makes Linux Kernel supply a GDB remote debug interface. Then GDB in current machine or remote machine can debug and trace Linux kernel and user space program through GDB tracepoint and some other functions without stopping the Linux Kernel.

And even if the board doesn't have GDB on it and doesn't have interface for remote debug. It can debug the Linux Kernel using offline debug (See [/sys/kernel/debug/gtpframe and offline debug](#)).

http://www.youtube.com/watch?v=7nfGAbNsEZY or http://www.tudou.com/programs/view/fPu_koiKo38/ is the video that introduced KGTP in **English**.

http://www.infoq.com/cn/presentations/gdb-sharp-knife-kgtp-linux-kernel is the video that introduced KGTP in **Chinese**.

KGTP supports **X86-32**, **X86-64**, **MIPS** and **ARM**.

KGTP supports most versions of Linux kernel (from **2.6.18** to **upstream**).

For new user of KGTP, please go to see [Quickstart](#).

Please go to [UPDATE](#) to get more info about KGTP update.

Get help or report issues

Please post issues to to https://github.com/teawater/kgtp/issues.

Or mail them to mailto:teawater@gmail.com?Subject=Report%20an%20issue%20of%20KGTP.

Or report it to QQ group **317654748**.

The KGTP team will try our best to help you.

Please goto https://code.google.com/p/kgtp/issues/list access the old issues list.

Table of different between GDB debug normal program and KGTP

This table is for the people that have experience using GDB debug normal program. It will help you understand and remember the function of KGTP.

Function	GDB debug normal program	GDB control KGTP debug Linux kernel
Preparatory work	Have a GDB installed in your system. Program built with "-g".	KGTP need GDB 7.6 or newer version because it use some new functions of GDB. If your system doesn't supply it, you can get new version GDB in http://code.google.com/p/gdbt/ and you can get an introduce about howto built new GDB step by step in there. You alse need do some preparatory work with Linux kernel and KGTP. Please goto Preparatory work before use KGTP get howto do it.
Attach	Use command "gdb -p pid" or GDB command "attach pid" can attach a program that running in the system.	Need insmod gtp.ko first, see Insmod the KGTP module. Then let GDB connect to KGTP, see Make GDB connect to gtp. Please note that after GDB connect to KGTP, Linux kernel will not stop.
Breakpoints	GDB command "b place_will_stop", let program execute after this command.	KGTP doesn't support breakpoints but it support tracepoints.

	Then programe will stop in the place that setup a breakpoint.	Tracepoints can be considered as a special kind of breakpoints. It can be setup in some place of Linux kernel and define some commands that you want to do in its actions. When tracepoints start, they will execute these commands when Linux kernel execute to these place. When tracepoint stop, you can use some GDB commands parse the data that get by tracepoints like what you do when program stop by breakpoints. **Difference** is breakpoints will stop the program But the tracepoints of KGTP not. Please goto GDB tracepoint get howto use it.
Memory read	After GDB stop the program(maybe doesn't need), it can read memory of program with GDB command "print", "x" and so on.	You can set special actions to collect memory to traceframe in tracepoints, and get the its value when tracepoint stop.collect expr1, expr2, ... Use tfind select the entry inside the trace frame info Or you can read memory directly when Linux kernel or program is running.Direct access the current value in normal mode
Step and continue	GDB can continue program execution with command "continue" and stop it with CTRL-C.	KGTP never stop the Linux kernel. But tracepoint can be start and stop.Start and stop the tracepoint

		Or use while-stepping tracepoint record Linux kernel with some times single step and Let KGTP switch to replay mode. Then it support execution commands (continue, step) and reverse-execute commands (reverse-continue, reverse-step). Use while-stepping let Linux kernel do single step
Backtrace	GDB can print backtrace of all stack frames with command "backtrace".	KGTP can do it too.Howto backtrace (stack dump)
Watchpoint	GDB can let programe stop when some memory access happen with watchpoint.	KGTP can record the memory access with watch tracepoint. Howto use watch tracepoint control hardware breakpoints to record memory access
Call function	GDB can call function of program with command "call function(xx,xx)".	KGTP can call function of Linux kernel with plugin.How to add plugin in C

Preparatory work before use KGTP

Linux kernel

If your system use the Linux kernel that is built by yourself

To use KGTP, your Linux kernel need open following options:

General setup --->
    [*] Kprobes

[*] Enable loadable module support --->

Kernel hacking --->
    [*] Debug Filesystem
    [*] Compile the kernel with debug info

Please rebuild your Linux kernel if you change any options of the config.

If use with Linux kernel of Android

The default Linux kernel config of Android should not support KGTP. To use KGTP, Linux kernel of Android need open following options:

[*] Enable loadable module support --->
General setup --->
    [*] Prompt for development and/or incomplete code/drivers
    [*] Kprobes
Kernel hacking --->
    [*] Debug Filesystem
    [*] Compile the kernel with debug info

Please rebuild your Linux kernel if you change any options of the Linux kernel config.

If your system use the Linux kernel from distribution

You need install some Linux kernel package.

Ubuntu

The standard method of install the Linux kernel debug image

1) Add debug source to the sources list of Ubuntu.

Create an /etc/apt/sources.list.d/ddebs.list by running the following line at a terminal:

> *echo "deb http://ddebs.ubuntu.com $(lsb_release -cs) main restricted universe multiverse" | *
> *sudo tee -a /etc/apt/sources.list.d/ddebs.list*

Stable releases (not alphas and betas) require three more lines adding to the same file, which is done by the following terminal command:

> *echo "deb http://ddebs.ubuntu.com $(lsb_release -cs)-updates main restricted universe multiverse*
> *deb http://ddebs.ubuntu.com $(lsb_release -cs)-security main restricted universe multiverse*
> *deb http://ddebs.ubuntu.com $(lsb_release -cs)-proposed main restricted universe multiverse" | *
> *sudo tee -a /etc/apt/sources.list.d/ddebs.list*

Import the debug symbol archive signing key:

> *sudo apt-key adv --keyserver keyserver.ubuntu.com --recv-keys 428D7C01*

Then run:

> *sudo apt-get update*

2) Get Linux kernel debug image

> *sudo apt-get install linux-image-$(uname -r)-dbgsym*

Then you can find Linux kernel debug image in "/usr/lib/debug/boot/vmlinux-$(uname -r)".

Please **note** that this step **Get Linux kernel debug image** need do again when Linux kernel update.

The second method of install the Linux kernel debug image

If you got some trouble with the standard method, please use following commands to install the Linux kernel debug image.

wget http://ddebs.ubuntu.com/pool/main/l/linux/linux-image-$(uname -r)-dbgsym_$(dpkg -s linux-image-$(uname -r) | grep ^Version: | sed 's/Version: //')_$(uname -i | sed 's/x86_64/amd64/').ddeb
sudo dpkg -i linux-image-$(uname -r)-dbgsym_$(dpkg -s linux-image-$(uname -r) | grep ^Version: | sed 's/Version: //')_$(uname -i | sed 's/x86_64/amd64/').ddeb

Please **note** that this method need do again when Linux kernel update.

Install the Linux kernel headers

sudo apt-get install linux-headers-generic

Install the Linux kernel source

New way

Install package that we need:

sudo apt-get install dpkg-dev

Get the Linux kernel source:

apt-get source linux-image-$(uname -r)

Then you can find Linux kernel directory in current directory.

Move this directory to "/build/buildd/".

Old way

Install the source package:

sudo apt-get install linux-source

Uncompress the source package:

sudo mkdir -p /build/buildd/
sudo tar vxjf /usr/src/linux-source-$(uname -r | sed 's/-.//').tar.bz2 -C /build/buildd/*
sudo rm -rf /build/buildd/linux-$(uname -r | sed 's/-.//')*
sudo mv /build/buildd/linux-source-$(uname -r | sed 's/-.*//') /build/buildd/linux-$(uname -r | sed 's/-.//')*

Please **note** that this step **Install the Linux kernel source** need do again when Linux kernel update.

Fedora

Install the Linux kernel debug image

Use following command:

sudo debuginfo-install kernel

Or:

sudo yum --enablerepo=fedora-debuginfo install kernel-debuginfo

Then you can find Linux kernel debug image in "/usr/lib/debug/lib/modules/$(uname -r)/vmlinux".

Install the Linux kernel devel package

sudo yum install kernel-devel-$(uname -r)

Please **note** that after update the Linux kernel package, you may need to call this command.

Make sure current Linux kernel debug image is right

GDB open the right Linux kernel debug image is an very important because GDB will get the debug info and address info from it. So before you use KGTP, please do the check to make sure about it.

There are 2 ways to do the check, what I suggest is do both of them to make sure Linux kernel debug image is right.

Please **note** that if you determine you use the right Linux kernel debug image, but cannot pass these ways. Please see <u>Handle the issue that Linux kernel debug image's address info is not same with Linux kernel when it running</u>.

Where is the current Linux kernel debug image

In UBUNTU, you can find it in "/usr/lib/debug/boot/vmlinux-$(uname -r)".

In Fedora, you can find it in "/usr/lib/debug/lib/modules/$(uname -r)/vmlinux".

If you build Linux kernel with yourself, file "vmlinux" in the Linux kernel build directory is the debug image.

Use /proc/kallsyms

In the system that its Linux kernel is what you want to trace, use following command to get the address of sys_read and sys_write:

```
sudo cat /proc/kallsyms | grep sys_read
ffffffff8117a520 T sys_read
sudo cat /proc/kallsyms | grep sys_write
ffffffff8117a5b0 T sys_write
```

Then we can get that the address of sys_read is 0xffffffff8117a520 and the address of sys_write is 0xffffffff8117a5b0.

After that use GDB get address of sys_read and sys_write from Linux kernel debug image:

```
gdb ./vmlinux
(gdb) p sys_read
$1 = {long int (unsigned int, char *, size_t)} 0xffffffff8117a520
<sys_read>
(gdb) p sys_write
$2 = {long int (unsigned int, const char *, size_t)} 0xffffffff8117a5b0
<sys_write>
```

The address of sys_read and sys_write is same, so the Linux kernel debug image is right.

Use linux_banner

```
sudo gdb ./vmlinux
(gdb) p linux_banner
$1 = "Linux version 3.4.0-rc4+ (teawater@teawater-
Precision-M4600) (gcc version 4.6.3 (GCC) ) #3 SMP Tue Apr
24 13:29:05 CST 2012\n"
```

This linux_banner is the kernel info inside the Linux kernel debug image.

After that, connect to KGTP following the way in Make GDB connect to gtp connect to KGTP and print linux_banner again.

```
(gdb) target remote /sys/kernel/debug/gtp
Remote debugging using /sys/kernel/debug/gtp
0x0000000000000000 in irq_stack_union ()
(gdb) p linux_banner
$2 = "Linux version 3.4.0-rc4+ (teawater@teawater-Precision-
M4600) (gcc version 4.6.3 (GCC) ) #3 SMP Tue Apr 24 13:29:05 CST
2012\n"
```

This linux_banner is the kernel info that Linux kernel that KGTP is tracing. If it is same with the prev kernel info, the Linux kernel debug image is right.

Handle the issue that Linux kernel debug image's address info is not same with Linux kernel when it running

In X86_32, you will found that the Linux kernel debug image's address info is not same with Linux kernel when it running through the ways in [Make sure current Linux kernel debug image is right](). And you determine the Linux kernel debug image is right.

This issue is because:

> *Processor type and features --->*
>   *(0x1000000) Physical address where the kernel is loaded*
>   *(0x100000) Alignment value to which kernel should be aligned*

The values of these two options are different. Please **note** that the "Physical address where the kernel is loaded" is not showed in config sometimes. You can get its value through search "PHYSICAL_START".

You can handle this issue through change "Alignment value to which kernel should be aligned" same with "Physical address where the kernel is loaded".

This issue doesn't affect X86_64.

Get KGTP

Get KGTP through http

Please goto https://github.com/teawater/kgtp/archive/master.zip get the upstream version of KGTP.

Please goto https://github.com/teawater/kgtp/archive/release.zip get the last release of KGTP.

Get KGTP through git

Following command will get the upstream version of KGTP:

git clone https://github.com/teawater/kgtp.git

Following command will get the last release version of KGTP:

git clone https://github.com/teawater/kgtp.git -b release

Mirrors

https://code.csdn.net/teawater/kgtp

https://www.gitshell.com/teawater/kgtp/

https://git.oschina.net/teawater/kgtp

Config KGTP

Following part is the default config of KGTP inside the Makefile. With this config, KGTP will build together with current kernel that running on this machine.

KERNELDIR := /lib/modules/`uname -r`/build
CROSS_COMPILE :=

KERELDIR is set to the directory which holds the kernel you want to build for. By default, it is set to the kernel that you are running.

Please note that this directory should be Linux kernel build directory or linux-headers directory but not the source directory but not the Linux kernel source directory. And the Linux kernel build directory should be used after build successful.

CROSS_COMPILE is set to the prefix name of compiler that you want to build KGTP. Empty to compile with your default compiler.

ARCH is the architecture.

Or you can choose which kernel you want build with and which compiler you want use by change Makefile.

For example:

KERNELDIR := /home/teawater/kernel/bamd64
CROSS_COMPILE :=x86_64-glibc_std-
ARCH := x86_64

KERNELDIR is set to /home/teawater/kernel/bamd64. Compiler will use x86_64-glibc_std-gcc.

Compile KGTP

Normal compile

cd kgtp/
make

In some build environment (for example Android) will get some error with user space program getmod or getframe. Please ignore this error and use the gtp.ko in this directory.

If you get error message "/usr/bin/ld: cannot find -lc" in Fedora, please use following command handle it.

sudo yum install glibc-static

Compile KGTP with some special config

Most of time, KGTP can auto select right options to build with Various versions of Linux kernel.

But if you want config special options with yourself, you can read following part:

With this option, KGTP will not auto select any build options.

make AUTO=0

With this option, KGTP will use simple frame instead of KGTP ring buffer.

The simple frame doesn't support gtpframe_pipe. It just for debug KGTP.

make AUTO=0 FRAME_SIMPLE=1

With this option, $clock will return rdtsc value instead of local_clock.

make AUTO=0 CLOCK_CYCLE=1

With this option, KGTP will use procfs instead of debugfs.

make AUTO=0 USE_PROC=1

The options can use together, for example:

make AUTO=0 FRAME_SIMPLE=1 CLOCK_CYCLE=1

Install and uninstall KGTP

KGTP don't need to be install because it can insmod directly inside its directory (See Insmod the KGTP module). But if you need, you can install it to your system.

Install:

cd kgtp/
sudo make install

Uninstall:

cd kgtp/
sudo make uninstall

Use KGTP with DKMS

You can use KGTP with DKMS if you want it.

Following commands will copy the files of KGTP to the directory that DKMS need.

cd kgtp/
sudo make dkms

Then you can use DKMS commands to control KGTP. Please goto http://linux.dell.com/dkms/manpage.html to see how to use DKMS.

Use KGTP patch for Linux kernel

Most of time, you don't need KGTP patch because KGTP can build as a LKM and very easy to use. But to help some people include KGTP to them special Linux Kernel tree, KGTP supply patches for Linux kernel.

In the KGTP directory:

- **gtp_3.7_to_upstream.patch** is the patch for Linux kernel from 3.7 to upstream.
- **gtp_3.0_to_3.6.patch is the patch** for Linux kernel from 3.0 to 3.6.
- **gtp_2.6.39.patch is the patch** for Linux kernel 2.6.39.
- **gtp_2.6.33_to_2.6.38.patch** is the patch for Linux kernel from 2.6.33 to 2.6.38.
- **gtp_2.6.20_to_2.6.32.patch** is the patch for Linux kernel from 2.6.20 to 2.6.32.
- **gtp_older_to_2.6.19.patch** is the patch for Linux kernel 2.6.19 and older version.

Install GDB for KGTP

The GDB that older than 7.6 have some bugs of tracepoint. And some functions of GDB are not very well.

So if your GDB is older than 7.6 please go to https://code.google.com/p/gdbt/ to get howto install GDB for KGTP. It supplies sources of UBUBTU, CentOS, Fedora, Mandriva, RHEL, SLE, openSUSE. Also have static binary for others.

If you have issue about GDB please get help according to Get help or report issues about KGTP.

How to let GDB connect to KGTP

To use KGTP function need let GDB connect to KGTP first.

Normal Linux

Insmod the KGTP module

If you have installed KGTP in your system, you can:

> *sudo modprobe gtp*

Or you can use the kgtp module in the directory.

> *cd kgtp/*
> *sudo insmod gtp.ko*

Handle the issue that cannot find "/sys/kernel/debug/gtp"

If you got this issue, please make sure "Debug Filesystem" is opened in your kernel config first. Please goto [If your system use the Linux kernel that is built by yourself](#) see how to open it.

If it is opened, please use following command mount sysfs.

sudo mount -t sysfs none /sys/

Maybe you will got some error for examle "sysfs is already mounted on /sys". Please ignore it.

please use following command mount debugfs.

mount -t debugfs none /sys/kernel/debug/

Then you can find "/sys/kernel/debug/gtp".

Make GDB connect to gtp

Load Linux kernel debug image to GDB

You can open GDB with following command to load image:

gdb kernel_debug_image_file

Or after you open GDB, use following GDB command to load image:

file kernel_debug_image_file

According to "Where is the current Linux kernel debug image", You can find Linux kernel debug image (kernel_debug_image_file).

Please **note** that let GDB open a right vmlinux file is very important for KGTP. Please goto "Make sure current Linux kernel debug image is right" get how to do it.

GDB on the current machine

sudo gdb ./vmlinux
(gdb) target remote /sys/kernel/debug/gtp
Remote debugging using /sys/kernel/debug/gtp
0x0000000000000000 in ?? ()

After that, you can begin to use GDB command trace and debug the Linux Kernel.

GDB on remote machine

Use nc map the KGTP interface to port 1024.

sudo su
nc -l 1234 </sys/kernel/debug/gtp >/sys/kernel/debug/gtp
#(nc -l -p 1234 </sys/kernel/debug/gtp >/sys/kernel/debug/gtp for
old version netcat.)

After that, nc will hang there to wait connection.

Let gdb connect to the port 1234.
gdb-release ./vmlinux
(gdb) target remote xxx.xxx.xxx.xxx:1234

After that, you can begin to use GDB command trace and debug the Linux Kernel.

Android

This video introduces use GDB connect to the KGTP in the Android, Please goto http://youtu.be/_UGN2j8Ctg0 or http://www.tudou.com/programs/view/FjkQ6HhPnfE/ to see it.

Insmod the KGTP module

First, make sure ADB has connected with Android.

Second, copy KGTP module to Android.

   sudo adb push gtp.ko /
Directory "/" may be read-only. You can choice other directory or use command "sudo adb shell mount -o rw,remount /" remount the directory to can write.

Third, insmod the module.

   adb shell insmod /gtp.ko

Handle the issue that cannot find "/sys/kernel/debug/gtp"

If you got this issue, please make sure "Debug Filesystem" is opened in your kernel config first. Please goto
If_your_system_use_the_Linux_kernel_that_is_built_by_yourself see howto "If use with Linux kernel of Android" see howto open it.

If it is opened, please use following command mount sysfs.

sudo adb shell mount -t sysfs none /sys/

Maybe you will got some error for examle "Device or resource busy". Please ignore it.

please use following command mount debugfs.

sudo adb shell mount -t debugfs none /sys/kernel/debug/

Then you can find "/sys/kernel/debug/gtp".

GDB connect to the KGTP

Use nc map the KGTP interface to port 1024.

adb forward tcp:1234 tcp:1234
adb shell "nc -l -p 1234 </sys/kernel/debug/gtp
>/sys/kernel/debug/gtp"
#(adb shell "nc -l 1234 </sys/kernel/debug/gtp
>/sys/kernel/debug/gtp" for new version netcat.)

After that, nc will hang there to wait connection.

Let gdb connect to the port 1234.

gdb-release ./vmlinux
(gdb) target remote :1234

After that, you can begin to use GDB command trace and debug the Linux Kernel.

Add module symbols to GDB

Sometimes you need to add a Linux kernel module's symbols to GDB to debug it.

Add symbols with hand is not very easy, so KGTP package include an GDB python script "getmod.py" and a program "getmod" can help you.

How to use getmod

"getmod" is written by C so you can use it anywhere even if in an embedded environment.

For example:

> *#Following command save Linux Kernel module info to the file*
> */tmp/mi in GDB*
> *#command format.*
> *sudo getmod >/tmp/mi*
> *#in gdb part:*
> *(gdb) source /tmp/mi*
> *add symbol table from file "/lib/modules/2.6.39-*
> *rc5+/kernel/fs/nls/nls_iso8859-1.ko" at*
> > *.text_addr = 0xf80de000*
> > *.note.gnu.build-id_addr = 0xf80de088*
> > *.exit.text_addr = 0xf80de074*
> > *.init.text_addr = 0xf8118000*
> > *.rodata.str1.1_addr = 0xf80de0ac*
> > *.rodata_addr = 0xf80de0c0*
> > *__mcount_loc_addr = 0xf80de9c0*
> > *.data_addr = 0xf80de9e0*
> > *.gnu.linkonce.this_module_addr = 0xf80dea00*
> *#After this GDB command, all the Linux Kernel module info is loaded*
> *into GDB.*

If you use remote debug or offline debug, maybe you need change the base directory. Following example is for it.

> *#/lib/modules/2.6.39-rc5+/kernel is replaced to sudo ./getmod -r*
> */home/teawater/kernel/b26*
> *sudo ./getmod -r /home/teawater/kernel/b26 >~/tmp/mi*

How to use getmod.py

Please **note** that static build GDB that download from https://code.google.com/p/gdbt/ cannot use getmod.py.

Connect to KGTP before use the getmod.py.

(gdb) source ~/kgtp/getmod.py

Then this script will auto load the Linux kernel module's symbols to GDB.

Howto use GDB control KGTP trace and debug Linux kernel

Direct access the current value in normal mode

After GDB connect to KGTP, if it doesn't select any a entry of trace frame bufffer with GDB command "tfind", GDB in the normal mode. Then you can direct access the current value of memory (Linux kernel or the user space program) and the trace state variables without stop anything.

If you have selected a trace frame entry, use GDB command "tfind -1" to return to normal mode. Please goto Use tfind select the entry inside the trace frame info info about GDB command "tfind".

The memory of Linux kernel

For example, you can access to "jiffies_64" with following command:

(gdb) p jiffies_64

Or you can access to the first entry of "static LIST_HEAD(modules)" with following command:

(gdb) p *((struct module *)((char *)modules->next - ((size_t) &(((struct module *)0)->list))))

Or you can access to the CPU0 memory info of "DEFINE_PER_CPU(struct device *, mce_device);":

p *(struct device *)(__per_cpu_offset[0]+(uint64_t)(&mce_device))

If you want show more than one variables with one GDB command, please use following example:

(gdb) printf "%4d %4d %4d %4d %4d %4d %18d %lu\n", this_rq->cpu, this_rq->nr_running, this_rq->nr_uninterruptible, nr_active, calc_load_tasks->counter, this_rq->calc_load_active, delta, this_rq-

>calc_load_update
2 1 0 0 0 0 673538312 717077240

the trace state variables

You can access value of TSV with the command that same with access memory.

Please goto [How to use trace state variables](#) get more info about TSV.

GDB tracepoint

Tracepoint is that GDB define some addresses and some actions and put them to the target (KGTP). After tracepoint start, , KGTP will do these actions (Some of them will collect data and save them to tracepoint frame buffer) when Linux kernel execution to there addresses. After that, Linux kernel will keep execution.

KGTP supply some interfaces that GDB or other programe can take the data of tracepoint frame buffer out to parse.

About these interfaces, this doc have introduced "/sys/kernel/debug/gtp". And will introduce "/sys/kernel/debug/gtpframe" and "/sys/kernel/debug/gtpframe_pipe" later.

Doc of GDB tracepoint in http://sourceware.org/gdb/current/onlinedocs/gdb/Tracepoints.html.

set tracepoint

The trace command is very similar to the break command. Its argument location can be a source line, a function name, or an address in the target program. The trace command defines a tracepoint, which is a address or some addresses that KGTP do some actions in it.

Here are some examples of using the trace command:

(gdb) trace foo.c:121   // a source file and line number

(gdb) trace +2   // 2 lines forward

(gdb) trace my_function // first source line of function

(gdb) trace *my_function // EXACT start address of function

(gdb) trace *0x2117c4   // an address

Howto handle the function is there but set tracepoint on it got fail

GCC will inline some static function to increase the performance. You cannot set tracepoint on the function name because object file doesn't have symbol of inline function.

You can use "trace filename:line" to set tracepoint on it.

How to set tracepoint condition

http://sourceware.org/gdb/current/onlinedocs/gdb/Tracepoint-Conditions.html

Like breakpoints, we can set conditions on tracepoints. The speed of tracepoints is faster than breakpoints because KGTP can do all the condition checks.

For example:

(gdb) trace handle_irq if (irq == 47)
This action of tracepoint 1 will work only when irq number is 47.

And you can use GDB command "condition" to specify the condition of a tracepoint. GDB command "condition N COND" will specify tracepoint number N to trace only if COND is true.

For example:

(gdb) trace handle_irq
(gdb) condition 1 (irq == 47)
GDB command "info tracepoint" will show the ID of the tracepoint.

Value of $bpnum is the last ID of GDB tracepoint. Then you can use GDB command "condtion" set the condition of last tracepoint without get its ID. For example:

(gdb) trace handle_irq
(gdb) condition $bpnum (irq == 47)

How to handle error "Unsupported operator (null) (52) in expression."

If you use condition about string, you will got this error when you call "tstart".

To handle it, you can convent the char to int to handle this issue, for example:

(gdb) p/x 'A'
$4 = 0x41
(gdb) condition 1 (buf[0] == 0x41)

actions [num]

This command will prompt for a list of actions to be taken when the tracepoint is hit. If the tracepoint number num is not specified, this command sets the actions for the one that was most recently defined (so that you can define a tracepoint and then say actions without bothering about its number). You specify the actions themselves on the following lines, one action at a time, and terminate the actions list with a line containing just end. So far, the only defined actions are collect, teval, and while-stepping.

collect expr1, expr2, ...

Collect values of the given expressions when the tracepoint is hit. This command accepts a comma-separated list of any valid expressions. In addition to global, static, or local variables, the following special arguments are supported:

> *$regs   Collect all registers.*
> *$args   Collect all function arguments.*
> *$locals Collect all local variables.*

Please **note** that collect an pointer (collect ptr) will just collect the address of this pointer. Add a * before ptr will make action collect the data that pointer point to(collect *ptr).

teval expr1, expr2, ...

Evaluate the given expressions when the tracepoint is hit. This command accepts a comma-separated list of expressions. The results are discarded, so this is mainly useful for assigning values to trace state variables (see Simple trace state variables) without adding those values to the trace buffer, as would be the case if the collect action were used.

while-stepping n

Please goto [Use while-stepping let Linux kernel do single step](#) see howto use it.

Start and stop the tracepoint

Tracepoint will exec actions only when it is starting use this GDB command:

(gdb) tstart

It will stop by this GDB command:

(gdb) tstop

Enable and disable the tracepoint

Like breakpoint, tracepoint can be control by GDB commands "enable" and "disable". But please **note** that it only useful when tracepoint stop.

Use tfind select the entry inside the trace frame info

GDB command "tfind" is used to select a entry of trace frame bufffer when tracepoint stop.

When GDB inside "tfind" mode, it will just show the values of this entry that the tracepoint action collect. So it will output some error when print some values that action doesn't collect for example the argument of function. That is not a bug, please don't worry about it.

Use "tfind" again will select next entry. "tfind id" will select entry id.

To return to normal mode([Direct access the current value in normal mode](#)), please use GDB command "tfind -1". Please goto http://sourceware.org/gdb/current/onlinedocs/gdb/tfind.html get more info about it.

How to handle error "No such file or directory."

When GDB cannot find the source code of Linux kernel, it will show this error message. For example:

> *(gdb) tfind*
> *Found trace frame 1, tracepoint 1*
> *#0 vfs_read (file=0xffff8801c36e6400, buf=0x7fff51a8f110*
> *<Address 0x7fff51a8f110 out of bounds>, count=16,*
> * pos=0xffff8801761dff48) at /build/buildd/linux-*
> *3.2.0/fs/read_write.c:365*
> *365   /build/buildd/linux-3.2.0/fs/read_write.c: No such file or*
> *directory.*

You can use GDB command "set substitute-path" to handle it. The prev example, the Linux kernel source is in "/build/buildd/test/linux-3.2.0/". But vmlinux let GDB find Linux kernel source in "/build/buildd/linux-3.2.0/". You can handle it with:

> *(gdb) set substitute-path /build/buildd/linux-3.2.0/*
> */build/buildd/test/linux-3.2.0/*
> *(gdb) tfind*
> *Found trace frame 1, tracepoint 1*
> *#0 vfs_read (file=0xffff8801c36e6400, buf=0x7fff51a8f110*
> *<Address 0x7fff51a8f110 out of bounds>, count=16,*
> * pos=0xffff8801761dff48) at /build/buildd/linux-*
> *3.2.0/fs/read_write.c:365*
> *365   {*

GDB have some other commands to handle the source code issue. Please goto http://sourceware.org/gdb/current/onlinedocs/gdb/Source-Path.html get the introduce about them.

Save the trace frame info to a file

/sys/kernel/debug/gtpframe supplies trace frame in tfile format (GDB can parse it) when KGTP is stop.

Please **note** that some "cp" cannot handle it very well, please use "cat /sys/kernel/debug/gtpframe > ./gtpframe" to copy it.

You can open file gtpframe when you want:

> *(gdb) target tfile ./gtpframe*
> *Tracepoint 1 at 0xffffffff8114f3dc: file /home/teawater/kernel/linux-2.6/fs/readdir.c, line 24.*
> *Created tracepoint 1 for target's tracepoint 1 at 0xffffffff8114f3c0.*
> *(gdb) tfind*
> *Found trace frame 0, tracepoint 1*
> *#0 vfs_readdir (file=0xffff880036e8f300, filler=0xffffffff8114f240 <filldir>, buf=0xffff880001e5bf38)*
> *   at /home/teawater/kernel/linux-2.6/fs/readdir.c:24*
> *24   {*

Show and save the tracepoint

You can use GDB command "info tracepoints" to show all the tracepoints.

You can use GDB command "save tracepoints filename" to save the commands that setup the tracepoints and actions into file filename. Then you use use GDB commands "source filename" to setup this tracepints again.

Delete tracepoint

GDB command "delete id" will delete tracepoint id. If "delete" without argument, it will delete all the tracepoint.

Use tracepoint get register info from a point of kernel

The following is an example that records the value of all registers when "vfs_readdir" is called.

(gdb) target remote /sys/kernel/debug/gtp
(gdb) trace vfs_readdir
Tracepoint 1 at 0xc01a1ac0: file
/home/teawater/kernel/linux-2.6/fs/readdir.c, line 23.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
>collect $reg
>end
(gdb) tstart
(gdb) shell ls
(gdb) tstop
(gdb) tfind
Found trace frame 0, tracepoint 1
#0 0xc01a1ac1 in vfs_readdir (file=0xc5528d00, filler=0xc01a1900 <filldir64>,
*   buf=0xc0d09f90) at /home/teawater/kernel/linux-2.6/fs/readdir.c:23*
23   /home/teawater/kernel/linux-2.6/fs/readdir.c: No such file or directory.
*   in /home/teawater/kernel/linux-2.6/fs/readdir.c*
(gdb) info reg
eax   0xc5528d00   -984445696
ecx   0xc0d09f90   -1060069488
edx   0xc01a1900   -1072031488
ebx   0xfffffff7   -9
esp   0xc0d09f8c   0xc0d09f8c
ebp   0x0   0x0
esi   0x8061480   134616192
edi   0xc5528d00   -984445696
eip   0xc01a1ac1   0xc01a1ac1 <vfs_readdir+1>
eflags   0x286   [PF SF IF]
cs   0x60   96
ss   0x8061480   134616192
ds   0x7b   123
es   0x7b   123
fs   0x0   0
gs   0x0   0
(gdb) tfind

Found trace frame 1, tracepoint 1
0xc01a1ac1   23   in /home/teawater/kernel/linux-2.6/fs/readdir.c
(gdb) info reg
eax   0xc5528d00   -984445696
ecx   0xc0d09f90   -1060069488
edx   0xc01a1900   -1072031488
ebx   0xfffffff7   -9
esp   0xc0d09f8c   0xc0d09f8c
ebp   0x0   0x0
esi   0x8061480   134616192
edi   0xc5528d00   -984445696
eip   0xc01a1ac1   0xc01a1ac1 <vfs_readdir+1>
eflags   0x286   [PF SF IF]
cs   0x60   96
ss   0x8061480   134616192
ds   0x7b   123
es   0x7b   123
fs   0x0   0
gs   0x0   0

Use tracepoint get the value of variable from a point of kernel

The following is an example that records the value of "jiffies_64" when the function "vfs_readdir" is called:

```
(gdb) target remote /sys/kernel/debug/gtp
(gdb) trace vfs_readdir
Tracepoint 1 at 0xc01ed740: file /home/teawater/kernel/linux-
2.6/fs/readdir.c, line 24.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
>collect jiffies_64
>collect file->f_path.dentry->d_iname
>end
(gdb) tstart
(gdb) shell ls
arch    drivers  include  kernel   mm          Module.symvers
security  System.map  virt
block   firmware init    lib      modules.builtin net          sound    t
vmlinux
crypto  fs       ipc     Makefile modules.order   scripts       source
usr        vmlinux.o
(gdb) tstop
(gdb) tfind
Found trace frame 0, tracepoint 1
#0  0xc01ed741 in vfs_readdir (file=0xf4063000, filler=0xc01ed580
<filldir64>, buf=0xd6dfdf90)
    at /home/teawater/kernel/linux-2.6/fs/readdir.c:24
24      {
(gdb) p jiffies_64
$1 = 4297248706
(gdb) p file->f_path.dentry->d_iname
$1 = "b26", '\000' <repeats 28 times>
```

Show all the traced data of current frame

After use "tfind" select an entry, you can use "tdump" to do it.

> *(gdb) tdump*
> *Data collected at tracepoint 1, trace frame 0:*
> *$cr = void*
> *file->f_path.dentry->d_iname =*
> *"gtp\000.google.chrome.g05ZYO\000\235\337\000\000\000\000\200*
> *067k\364\200\067", <incomplete sequence \364>*
> *jiffies_64 = 4319751455*

Get status of tracepoint

Please use GDB command "tstatus".

Set the trace buffer into a circular buffer

http://sourceware.org/gdb/current/onlinedocs/gdb/Starting-and-Stopping-Trace-Experiments.html

The frame buffer is not a circular buffer by default. When the buffer is full, the tracepoint will stop.

Following command will set frame buffer to a circular buffer. When the buffer is full, it will auto discard traceframes (oldest first) and keep trace.

(gdb) set circular-trace-buffer on

Do not stop tracepoint when the GDB disconnects

http://sourceware.org/gdb/current/onlinedocs/gdb/Starting-and-Stopping-Trace-Experiments.html

KGTP will stop tracepoint and delete the trace frame when GDB disconnects with it by default.

Following command will open the KGTP disconnect-trace. After that, when GDB disconnects with KGTP, KGTP will not stop tracepoint. And after GDB reconnects to KGTP, it can keep control of KGTP like nothing happened.

> *(gdb) set disconnected-tracing on*

kprobes-optimization and the execution speed of tracepoint

The tracepoint is execution together with Linux kernel. So it speed will affect the speed the system.

The KGTP tracepoint base on Linux kernel kprobe. Because the normal kprobe base on breakpoint instruction, so it is not very fast.

But if arch of kernel is X86_64 or X86_32 and kernel config didn't open "Preemptible Kernel" (PREEMPT), the kprobe is speed up by kprobes-optimization (CONFIG_OPTPROBES) that make kprobe very fast.

To make sure about that, you can use following command in terminal:

> *sysctl -A | grep kprobe*
> *debug.kprobes-optimization = 1*

That means that your kernel support kprobes-optimization.

Please **note** that some KGTP functions will make this tracepoint use simple kprobe even if this Kernel support kprobes-optimization. This doc will add note when introduce these functions. Please avoid using them when you really care about the tracepoint speed.

How to use trace state variables

http://sourceware.org/gdb/current/onlinedocs/gdb/Trace-State-Variables.html

Trace state variable is referred to as the TSV.

TSV can be accessed in tracepoint action and condition or direct access by GDB command.

Please **note** that just GDB 7.2.1 and later versions support use trace state variables directly, the old version of GDB can show the value of trace state variables through command "info tvariables".

Simple trace state variables

Define a trace state variable $c.

(gdb) tvariable $c

Trace state variable $c is created with initial value 0. The following action uses $c to count how many irqs happened in the kernel.

(gdb) target remote /sys/kernel/debug/gtp
(gdb) trace handle_irq
(gdb) actions
Enter actions for tracepoint 3, one per line.
End with a line saying just "end".
>collect $c   #Save current value of $c to the trace frame buffer.
>teval $c=$c+1 #Increase the $c.
>end

Also, you can set a value of variable to trace state variable, but don't forget covert variable to "uint64_t".

>teval $c=(uint64_t)a

You can get the current value of $c while the trace is running or stopped.

(gdb) tstart
(gdb) info tvariables
$c   0   31554
(gdb) p $c
$5 = 33652
(gdb) tstop
(gdb) p $c
$9 = 105559

When using tfind, you can parse the trace frame buffer. If the value of a trace state variable is collected, you can parse it out.

(gdb) tstop
(gdb) tfind
(gdb) info tvariables
$c   0   0
(gdb) p $c
$6 = 0
(gdb) tfind 100
(gdb) p $c
$7 = 100

If need, the tracepoint action that access the simple trace state variables will auto lock a spin lock for trace state variables. So it can handle race condition issue about trace state variables.

The following example is OK even if it running a machine that have more than one CPU.

>*teval $c=$c+1*

Per_cpu trace state variables

Per_cpu trace state variables are special simple trace state variables.

When tracepoint action access to it, it will access to this CPU special trace state variables.

It have 2 advantages:

1. The tracepoint actions that access to per_cpu trace state variables don't have the race conditon issue. So it don't need lock the spin lock for trace state variables. It is faster than simple trace state variables on multi-core machine.

2. Write the action that count some CPU special thing with it is easier than simple trace state variables.

How to define

Per_cpu trace state variables have two types:

Local CPU variables

> *"per_cpu_"+string*

or

> *"p_"+string*

For example:

> *(gdb) tvariable $p_count*

When access this trace state variable in tracepoint actions, it will return the variable's value of CPU that this tracepoint actions running on.

CPU id variables

> *"per_cpu_"+string+CPU_id*

or

> *"p_"+string+CPU_id*

For example:

> *(gdb) tvariable $p_count0*
> *(gdb) tvariable $p_count1*
> *(gdb) tvariable $p_count2*
> *(gdb) tvariable $p_count3*

When access this trace state variable in tracepoint actions or GDB command line, it will return the variable's value of CPU CPU_id.

Follow example can auto define a CPU id variables for each CPU of this machine. (Please note that need let GDB connect to KGTP before use these commands.)

> *(gdb) set $tmp=0*
> *(gdb) while $tmp<$cpu_number*
> *>eval "tvariable $p_count%d",$tmp*
> *>set $tmp=$tmp+1*
> *>end*

Example 1

This example define a tracepoint that count the times that call vfs_read of each CPU.

```
tvariable $p_count
set $tmp=0
while $tmp<$cpu_number
  eval "tvariable $p_count%d",$tmp
  set $tmp=$tmp+1
  end
trace vfs_read
actions
  teval $p_count=$p_count+1
  end
```

Then you can show how many vfs_read in each CPU after "tstart":

```
(gdb) p $p_count0
$3 = 44802
(gdb) p $p_count1
$4 = 55272
(gdb) p $p_count2
$5 = 102085
(gdb) p $p_count3
```

Example 2

This example record stack dump of the function that close IRQ longest time of each CPU.

```
set pagination off

tvariable $bt=1024
tvariable $p_count
tvariable $p_cc
set $tmp=0
while $tmp<$cpu_number
eval "tvariable $p_cc%d",$tmp
set $tmp=$tmp+1
end

tvariable $ignore_error=1

trace arch_local_irq_disable
  commands
    teval $p_count=$clock
  end
trace arch_local_irq_enable if ($p_count && $p_cc < $clock -
$p_count)
  commands
    teval $p_cc = $clock - $p_count
    collect $bt
    collect $p_cc
    teval $p_count=0
  end

enable
set pagination on
```

Special trace state variables $current_task, $current_task_pid, $current_thread_info, $cpu_id, $dump_stack, $printk_level, $printk_format, $printk_tmp ,$clock, $hardirq_count, $softirq_count and $irq_count

KGTP special trace state variables $current_task, $current_thread_info, $cpu_id and $clock can very easy to access to some special value. You can see them when GDB connects to the KGTP. You can use them in tracepoint conditions or actions.

Access $current_task in tracepoint condition and action will get that returns of get_current().

Access $current_task_pid in tracepoint condition and action will get that returns of get_current()->pid.

Access $current_thread_info in tracepoint condition and action will get that returns of current_thread_info().

Access $cpu_id in tracepoint condition and action will get that returns of smp_processor_id().

Access $clock in tracepoint condition and action will get that returns of local_clock() that return the timestamp in nanoseconds.

$rdtsc is only available on X86 and X86_64 architecture. Access it in anytime will get current value of TSC with instruction RDTSC.

Access $hardirq_count in tracepoint condition and action will get that returns of hardirq_count().

Access $softirq_count in tracepoint condition and action will get that returns of softirq_count().

Access $irq_count in tracepoint condition and action will get that returns of irq_count().

And KGTP has other special trace state variables $dump_stack, $printk_level, $printk_format and $printk_tmp. All of them output their values directly, as can be seen in [Howto let tracepoint output value directly](#).

The following example counts in $c how many vfs_read calls that process 16663 does and collects the struct thread_info of current task:

> *(gdb) target remote /sys/kernel/debug/gtp*
> *(gdb) trace vfs_read if (((struct task_struct *)$current_task)->pid ==*
> *16663)*
> *(gdb) tvariable $c*

```
(gdb) actions
Enter actions for tracepoint 4, one per line.
End with a line saying just "end".
>teval $c=$c+1
>collect (*(struct thread_info *)$current_thread_info)
>end
(gdb) tstart
(gdb) info tvariables
Name          Initial    Current
$c            0          184
$current_task   0          <unknown>
$current_thread_info 0       <unknown>
$cpu_id       0          <unknown>
(gdb) tstop
(gdb) tfind
(gdb) p *(struct thread_info *)$current_thread_info
$10 = {task = 0xf0ac6580, exec_domain = 0xc07b1400, flags = 0,
status = 0, cpu = 1, preempt_count = 2, addr_limit = {
    seg = 4294967295}, restart_block = {fn = 0xc0159fb0
<do_no_restart_syscall>, {{arg0 = 138300720, arg1 = 11,
     arg2 = 1, arg3 = 78}, futex = {uaddr = 0x83e4d30, val = 11,
flags = 1, bitset = 78, time = 977063750,
     uaddr2 = 0x0}, nanosleep = {index = 138300720, rmtp = 0xb,
expires = 335007449089}, poll = {
     ufds = 0x83e4d30, nfds = 11, has_timeout = 1, tv_sec = 78,
tv_nsec = 977063750}}},
  sysenter_return = 0xb77ce424, previous_esp = 0, supervisor_stack
= 0xef340044 "", uaccess_err = 0}
```

Another example shows how much sys_read() executes in each CPU.

```
(gdb) tvariable $c0
(gdb) tvariable $c1
(gdb) trace sys_read
(gdb) condition $bpnum ($cpu_id == 0)
(gdb) actions
>teval $c0=$c0+1
>end
(gdb) trace sys_read
(gdb) condition $bpnum ($cpu_id == 1)
(gdb) actions
>teval $c1=$c1+1
>end
(gdb) info tvariables
Name          Initial    Current
$current_task   0          <unknown>
$cpu_id       0          <unknown>
$c0           0          3255
```

$c1   0   1904

sys_read() execute 3255 times in cpu0 and 1904 times in cpu1. Please note that this example just to howto use $cpu_id. Actually, this example use per_cpu trace state variables is better.

Special trace state variable $self_trace

$self_trace is different with the special trace state variables in the previous section. It is used to control the behavior of tracepoint.

In default, when tracepoint is triggered, the actions will not execute if the current_task is the a KGTP self process (GDB, netcat, getframe or some others process that access to the interface of KGTP).

If you want tracepoint actions execute with any task, please include a command access to the $self_trace in the actions i.e. add following command to the actions:

>*teval $self_trace=0*

Trace the function return with $kret

Sometime, set the tracepoint to the end of function is hard because the Kernel is compiled with optimization. At this time, you can get help from $kret.

$kret is a special trace state variable like $self_trace. When you set value of it inside the action of tracepoint, this tracepoint be set with kretprobe instead of kprobe. Then it can trace the end of this function.

Please note that this tracepoint must set in the first address of the function in format "**function_name**".

Following part is an example:

> *#"*(function_name)" format can make certain that GDB send the*
> *first address of function to KGTP.*
> *(gdb) trace *vfs_read*
> *(gdb) actions*
> *>teval $kret=0*
> *#Following part you can set commands that you want.*

Use $ignore_error and $last_errno to ignore the error of tstart

If KGTP got any error of tstart, this command will get fail.

But sometime we need ignore this error and let KGTP keep work. For example: If you set tracepoint on the inline function spin_lock. This tracepoint will be set to a lot of addresses that some of them cannot be set kprobe. It will make tstart get fail. You can use "$ignore_error" ignore this error.

And the last error number will available in "$last_errno".

(gdb) tvariable $ignore_error=1
This command will open ignore.

(gdb) tvariable $ignore_error=0
This command will close ignore.

Use $cooked_clock and $cooked_rdtsc the time without KGTP used

Access these two trace state variables can get the time without KGTP used. Then we can get more close to really time that a part of code used even if the actions of tracepoint is very complex. They will be introduce in Cookbook (coming soon).

Use $xtime_sec and $xtime_nsec get the timespec

Access these two trace state variables will return the time of day in a timespec that use getnstimeofday.

$xtime_sec will access to the second part of a timespec.

$xtime_nsec will access to the nanosecond part of a timespec.

Howto backtrace (stack dump)

Each time your program performs a function call, information about the call is generated. That information includes the location of the call in your program, the arguments of the call, and the local variables of the function being called. The information is saved in a block of data called a stack frame. The stack frames are allocated in a region of memory called the call stack.

Collect stack with $bt and use GDB command "backtrace"

Because this way is faster (just collect the stack when trace) and parse out most of info inside the call stack (it can show all the stack info that I introduce). So I suggest you use this way to do the stack dump.

First we need add the collect the stack command to the tracepoint action.

~~The general collect the stack command in GDB tracepoint is: In x86_32, following command will collect 512 bytes of stack.~~

~~>collect *(unsigned char *)$esp@512~~

~~In x86_64, following command will collect 512 bytes of stack.~~

~~>collect *(unsigned char *)$rsp@512~~

~~In MIPS or ARM, following command will collect 512 bytes of stack.~~

~~>collect *(unsigned char *)$sp@512~~

~~These commands is so hard to remember, and the different arch need different command.~~

KGTP have an special tracepoint trace state variable $bt. If tracepoint action access it, KGTP will auto collect the $bt size (default value is 512) stack. For example, this command will collect 512 bytes stack memory:

> *>collect $bt*

If you want to change size of $bt, you can use following GDB command before "tstart":

> *(gdb) tvariable $bt=1024*

Following part is an example about howto collect stack and howto use GDB parse it:

> *(gdb) target remote /sys/kernel/debug/gtp*
> *(gdb) trace vfs_readdir*
> *Tracepoint 1 at 0xffffffff8118c300: file*
> */home/teawater/kernel2/linux/fs/readdir.c, line 24.*
> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*
> *End with a line saying just "end".*
> *>collect $bt*
> *>end*
> *(gdb) tstart*
> *(gdb) shell ls*
> *1   crypto   fs   include kernel   mm*
> *Module.symvers security System.map vmlinux*
> *arch drivers hotcode.html init   lib   modules.builtin net*

```
sound     usr       vmlinux.o
block  firmware  hotcode.html~  ipc     Makefile  modules.order
scripts        source     virt
(gdb) tstop
(gdb) tfind
Found trace frame 0, tracepoint 1
#0  vfs_readdir (file=0xffff8800c5556d00, filler=0xffffffff8118c4b0
<filldir>, buf=0xffff880108709f40)
    at /home/teawater/kernel2/linux/fs/readdir.c:24
24     {
(gdb) bt
#0  vfs_readdir (file=0xffff8800c5556d00, filler=0xffffffff8118c4b0
<filldir>, buf=0xffff880108709f40)
    at /home/teawater/kernel2/linux/fs/readdir.c:24
#1  0xffffffff8118c689 in sys_getdents (fd=<optimized out>,
dirent=0x1398c58, count=32768) at
/home/teawater/kernel2/linux/fs/readdir.c:214
#2  <signal handler called>
#3  0x00007f00253848a5 in ?? ()
#4  0x00003efd32cddfc9 in ?? ()
#5  0x00002c15b7d04101 in ?? ()
#6  0x000019c0c5704bf1 in ?? ()
#7  0x0000000900000000 in ?? ()
#8  0x000009988cc8d269 in ?? ()
#9  0x000009988cc9b8d1 in ?? ()
#10 0x0000000000000000 in ?? ()
(gdb) up
#1  0xffffffff8118c689 in sys_getdents (fd=<optimized out>,
dirent=0x1398c58, count=32768) at
/home/teawater/kernel2/linux/fs/readdir.c:214
214            error = vfs_readdir(file, filldir, &buf);
(gdb) p buf
$1 = {current_dir = 0x1398c58, previous = 0x0, count = 32768,
error = 0}
(gdb) p error
$3 = -9
(gdb) frame 0
#0  vfs_readdir (file=0xffff8800c5556d00, filler=0xffffffff8118c4b0
<filldir>, buf=0xffff880108709f40)
    at /home/teawater/kernel2/linux/fs/readdir.c:24
24     {
```

From this example, we can see some GDB commands that parse the the call stack:

- **bt** is the alias of GDB commands backtrace that print a backtrace of the entire stack: one line per frame for all frames in the stack.

- **up n** is move n frames up the stack. For positive numbers n, this

advances toward the outermost frame, to higher frame numbers, to frames that have existed longer. n defaults to one.

- **down n** is move n frames down the stack. For positive numbers n, this advances toward the innermost frame, to lower frame numbers, to frames that were created more recently. n defaults to one. You may abbreviate down as do.

- **frame n** is select frame number n. Recall that frame zero is the innermost (currently executing) frame, frame one is the frame that called the innermost one, and so on. The highest-numbered frame is the one for main.

You can see that when you use up, down or frame to the different calll stack frame, you can output the value of the arguments and local variables of different call stack frame.

To get the more info about howto use GDB parse the call stack, please see http://sourceware.org/gdb/current/onlinedocs/gdb/Stack.html

Collect stack of current function's caller with $_ret

If you just want to collect stack of current function's caller, please use $_ret.

Please **note** that set the tracepoint that collect $_ret cannot in the first address of function.

For example:

(gdb) list vfs_read
360   }
361
362   EXPORT_SYMBOL(do_sync_read);
363
*364   ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)*
365   {
366   ssize_t ret;
367
368   if (!(file->f_mode & FMODE_READ))
369   return -EBADF;
(gdb) trace 368
Tracepoint 2 at 0xffffffff8117a244: file /home/teawater/kernel2/linux/fs/read_write.c, line 368.
(gdb) actions
Enter actions for tracepoint 2, one per line.
End with a line saying just "end".
>collect $_ret
>end
(gdb) tstart
(gdb) tstop
(gdb) tfind
Found trace frame 0, tracepoint 2
#0 vfs_read (file=0xffff880141c46000, buf=0x359bda0 <Address 0x359bda0 out of bounds>, count=8192, pos=0xffff88012fa49f48)
*   at /home/teawater/kernel2/linux/fs/read_write.c:368*
368   if (!(file->f_mode & FMODE_READ))
(gdb) bt
#0 vfs_read (file=0xffff880141c46000, buf=0x359bda0 <Address 0x359bda0 out of bounds>, count=8192, pos=0xffff88012fa49f48)
*   at /home/teawater/kernel2/linux/fs/read_write.c:368*
#1 0xffffffff8117a3ea in sys_read (fd=<optimized out>, buf=<unavailable>, count=<unavailable>)
*   at /home/teawater/kernel2/linux/fs/read_write.c:469*
Backtrace stopped: not enough registers or memory available to unwind further
(gdb) up

#1 0xffffffff8117a3ea in sys_read (fd=<optimized out>,
buf=<unavailable>, count=<unavailable>)
*   at /home/teawater/kernel2/linux/fs/read_write.c:469*
469   ret = vfs_read(file, buf, count, &pos);
(gdb) p ret
$2 = -9

You see that the caller of function vfs_read is sys_read. And the local variable ret of sys_read is -9.

Use $dump_stack to output stack dump through printk

Because this way need parse the stack when tracing and call printk inside, so it will be slow, unsafe, unclear and cannot access a lot of info of call stack. So I suggest you use the prev way to do stack dump.

KGTP has special trace state variable $dump_stack, "collect" it will let Linux Kernel output stack dump through printk.

Following example lets Linux Kernel show the stack dump of vfs_readdir:

target remote /sys/kernel/debug/gtp
trace vfs_readdir
  commands
    collect $dump_stack
  end

Then your kernel will printk like:

[22779.208064] gtp 1:Pid: 441, comm: python Not tainted 2.6.39-rc3+ #46
[22779.208068] Call Trace:
[22779.208072] [<fe653cca>] gtp_get_var+0x4a/0xa0 [gtp]
[22779.208076] [<fe653d79>] gtp_collect_var+0x59/0xa0 [gtp]
[22779.208080] [<fe655974>] gtp_action_x+0x1bb4/0x1dc0 [gtp]
[22779.208084] [<c05b6408>] ? _raw_spin_unlock+0x18/0x40
[22779.208088] [<c023f152>] ? __find_get_block_slow+0xd2/0x160
[22779.208091] [<c01a8c56>] ? delayacct_end+0x96/0xb0
[22779.208100] [<c023f404>] ? __find_get_block+0x84/0x1d0
[22779.208103] [<c05b6408>] ? _raw_spin_unlock+0x18/0x40
[22779.208106] [<c02e0838>] ? find_revoke_record+0xa8/0xc0
[22779.208109] [<c02e0c45>] ?
jbd2_journal_cancel_revoke+0xd5/0xe0
[22779.208112] [<c02db51f>] ?
__jbd2_journal_temp_unlink_buffer+0x2f/0x110
[22779.208115] [<fe655c4c>] gtp_kp_pre_handler+0xcc/0x1c0
[gtp]
[22779.208118] [<c05b8a88>]
kprobe_exceptions_notify+0x3d8/0x440
[22779.208121] [<c05b7d54>] ?
hw_breakpoint_exceptions_notify+0x14/0x180
[22779.208124] [<c05b95eb>] ? sub_preempt_count+0x7b/0xb0
[22779.208126] [<c0227ac5>] ? vfs_readdir+0x15/0xb0
[22779.208128] [<c0227ac4>] ? vfs_readdir+0x14/0xb0
[22779.208131] [<c05b9743>] notifier_call_chain+0x43/0x60
[22779.208134] [<c05b9798>]
__atomic_notifier_call_chain+0x38/0x50
[22779.208137] [<c05b97cf>] atomic_notifier_call_chain+0x1f/0x30

[22779.208140] [<c05b980d>] notify_die+0x2d/0x30
[22779.208142] [<c05b71c5>] do_int3+0x35/0xa0

Howto let tracepoint output value directly

In the previous parts, you may understand that to get a value from Linux kernel, you need to use a tracepoint "collect" action to save the value to the tracepoint frame and use the GDB command "tfind" to parse the value from the frame data.

But we want get the value directly sometimes, so KGTP supports two ways to output values directly.

Switch collect to output the value directly

KGTP has special trace state variables $printk_level, $printk_format and $printk_tmp to support this function.

$printk_level: if its value is 8 (this is the default value), "collect" action will save value to the tracepoint frame in the simple behavior.

If its value is 0-7, "collect" will output the value through "printk" directly, and value will be the level of printk. The level is:

0   KERN_EMERG   system is unusable
1   KERN_ALERT   action must be taken immediately
2   KERN_CRIT   critical conditions
3   KERN_ERR   error conditions
4   KERN_WARNING   warning conditions
5   KERN_NOTICE   normal but significant condition
6   KERN_INFO   informational
7   KERN_DEBUG   debug-level messages

$printk_format, collect printk will output value in the format that is set by it. The format is:

0   This is the default value.
*   If the size of collect value is 1, 2, 4 or 8, it will be output as an unsigned decimal.*
*   If not, it will be output as a hexadecimal string.*
1   Output value in signed decimal.
2   Output value in unsigned decimal.
3   Output value in unsigned hexadecimal.
4   Output value as a string.
5   Output value as a hexadecimal string.

$printk_tmp, to output the value of global variable need set to it first.

Following example shows a count number, pid, jiffies_64 and the file name

that call vfs_readdir:

```
(gdb) target remote /sys/kernel/debug/gtp
(gdb) tvariable $c
(gdb) trace vfs_readdir
(gdb) actions
>teval $printk_level=0
>collect $c=$c+1
>collect ((struct task_struct *)$current_task)->pid
>collect $printk_tmp=jiffies_64
>teval $printk_format=4
>collect file->f_path.dentry->d_iname
>end
```

Then your kernel will printk like:

```
gtp 1:$c=$c+1=41
gtp 1:((struct task_struct *)$current_task)->pid=12085
gtp 1:$printk_tmp=jiffies_64=4322021438
gtp 1:file->f_path.dentry->d_iname=b26
gtp 1:$c=$c+1=42
gtp 1:((struct task_struct *)$current_task)->pid=12085
gtp 1:$printk_tmp=jiffies_64=4322021438
gtp 1:file->f_path.dentry->d_iname=b26
```
"gtp 1" means that it was output by tracepoint 1.

Howto use watch tracepoint control hardware breakpoints to record memory access

Watch tracepoint can control hardware breakpoints to record the memory access through set some special trace state variables in its action.

Please **note** that watch tracepoint is just support by X86 and X86_64 now. And dynamic watch tracepoint just can work OK in Linux 2.6.27 and newer version because Linux 2.6.26 and older version have some IPI issues on smp support.

Trace state variables of watch tracepoint

Name	Written by normal tracepoint	Read by normal tracepoint	Written by static watch tracepoint	Read by static watch tracepoint	Written by dynamic watch tracepoint	Read by dynamic watch tracepoint
$watch_static	Not support	Not support	If "teval $watch_static=1", then this tracepoint is static watch tracepoint.	Not support	If "teval $watch_static=1", then this tracepoint is static watch tracepoint.	Not support
$watch_set_id	When this tracepoint want to setup a dynamic watch tracepoint, set a id of a dynamic watch tracepoint to $watch_set_id to point out which dynamic watch tracepoint you wan to setup.	Not support	Not support	Not support	Not support	Not support
$watch_set_addr	When this tracepoint want to setup a dynamic watch tracepoint, set the address of a dynamic watch tracepoint to $watch_set_addr to point out which dynamic watch tracepoint you wan to setup.	Not support	Not support	Not support	Not support	Not support
$watch_type	When this tracepoint want to setup a dynamic watch tracepoint, set the watch type of this dynamic watch tracepoint to $watch_type. 0 is exec. 1 is write. 2 is read or write.	Get the value that this tracepoint set to $watch_type	Set the type of this watch tracepoint.	Get the type of this watch tracepoint.	Set the default type of this watch tracepoint.	Get the type of this watch tracepoint when it really exec.
$watch_size	When this tracepoint want to setup a dynamic watch tracepoint, set the watch size of this dynamic watch tracepoint to $watch_size. The size should be 1, 2, 4, 8.	Get the value that this tracepoint set to $watch_size.	Set the size of this watch tracepoint.	Get the size of this watch tracepoint.	Set the default size of this watch tracepoint.	Get the size of this watch tracepoint when it really exec.
$watch_start	Set the address to a dynamic watch tracepoint(set by $watch_set_addr or $watch_set_id) and let it try to start	Get the result of this start. (It will fail becasue X86 just have 4 hardware	Not support	Not support	Not support	Not support

	work.	breakpoints.) Get 0 if success. If < 0 is the error id.	94			
$watch_stop	Set a address to $watch_stop will let a dynamic watch tracepoint that watch in this address stop.	Get the result of this stop.	Not support	Not support	Not support	Not support
$watch_trace_num	Not support	Not support	Not support	Not support	Not support	The tracepoint number that setup this dynamic watch tracepoint.
$watch_trace_addr	Not support	Not support	Not support	Not support	Not support	The tracepoint address that setup this dynamic watch tracepoint.
$watch_addr	Not support	Not support	Not support	The address that this watch tracepoint is watching.	Not support	The address that this watch tracepoint is watching.
$watch_val	Not support	Not support	Not support	The current value of the memory that this watch tracepoint is watching.	Not support	The current value of the memory that this watch tracepoint is watching.
$watch_prev_val	Not support	Not support	Not support	The previous value of the memory that this watch tracepoint is watching.	Not support	The previous value of the memory that this watch tracepoint is watching.
$watch_count	Not support	Not support	Not support	Not support	Not support	A special count for this watch tracepoint session.

Static watch tracepoint

You can use static watch tracepoint when you want watch value of a global variable or some memory that you can get its address directly. Following example is watch jiffies_64's write:

```
#Static watch tracepoint get watch address from tracepoint address.
trace *&jiffies_64
  actions
    #Set this watch tracepoint to static
    teval $watch_static=1
    #Watch memory write
    teval $watch_type=1
    teval $watch_size=8
    collect $watch_val
    collect $watch_prev_val
    collect $bt
  end
```

Dynamic watch tracepoint

If you want to watch value of a local variable or some memory that you just get get its address inside the function, you can use dynamic watch tracepoint. Following example is watch write of f->f_pos and f->f_op inside function get_empty_filp:

```
trace *1
  commands
    teval $watch_static=0
    teval $watch_type=1
    teval $watch_size=8
    collect $bt
    collect $watch_addr
    collect $watch_val
    collect $watch_prev_val
  end
```

Define a dynamic watch tracepoint. The address "1" of it is not the address of memory that it will watch. It just help tracepoint that setup this dynamic watch tracepoint can find it.

```
list get_empty_filp
trace 133
  commands
    teval $watch_set_addr=1
    teval $watch_size=4
    teval $watch_start=&(f->f_pos)
    teval $watch_size=8
    teval $watch_start=&(f->f_op)
  end
```

Define a normal tracepoint that start to watch f->f_pos and f->f_op inside function get_empty_filp.

```
trace file_sb_list_del
  commands
    teval $watch_stop=&(file->f_pos)
    teval $watch_stop=&(file->f_op)
  end
```

Define a normal tracepoint that stop the tracepoint that watch file->f_pos and file->f_op.

Use while-stepping let Linux kernel do single step

Please **note** that while-stepping is just support by X86 and X86_64 now.

Video about howto use while-stepping http://www.codepark.us/a/13.

Howto use while-stepping

while-stepping is a special tracepoint action that include some actions with it.

When tracepoints that its actions include "while-stepping n" execute, it will do n times single steps and executes the actions of while-stepping. For example:

```
trace vfs_read
#Because single step will make system slow, so use passcount or
condition to limit the execution times of tracepoint is better.
passcount 1
  commands
    collect $bt
    collect $step_count
    #do 2000 times single steps.
    while-stepping 2000
      #Following part is actions of "while-stepping 2000".
      #Because step maybe execute to other functions, so does not
access local variables is better.
      collect $bt
      collect $step_count
    end
  end
```

Please **note** that tracepoint will disable the interrupt of current CPU when it do single step. Access **$step_count** in actions will get the count of this step that begin with 1.

Read the traceframe of while-stepping

The data of different step that is recorded by while-stepping actions will be saved in different traceframe that you can use tfind ([Use tfind select the entry inside the trace frame info](#)) to select them.

Or you can switch KGTP to replay mode to select all the traceframe of a while-stepping tracepoint with GDB execution and reverse-execution commands. For example:

Use tfind select one the traceframe of a while-stepping tracepoint.

(gdb) tfind
Found trace frame 0, tracepoint 1
#0 vfs_read (file=0xffff8801f7bd4c00, buf=0x7fff74e4edb0
<Address 0x7fff74e4edb0 out of bounds>, count=16,
*   pos=0xffff8801f4b45f48) at /build/buildd/linux-*
3.2.0/fs/read_write.c:365
365   {

Following commands will swith KGTP to replay mode.

(gdb) monitor replay
(gdb) tfind -1
No longer looking at any trace frame
#0 vfs_read (file=0xffff8801f7bd4c00, buf=0x7fff74e4edb0
<Address 0x7fff74e4edb0 out of bounds>, count=16,
*   pos=0xffff8801f4b45f48) at /build/buildd/linux-*
3.2.0/fs/read_write.c:365
365   {

Then you can use execution commands.

(gdb) n
368   if (!(file->f_mode & FMODE_READ))
(gdb) p file->f_mode
$5 = 3

Set breakpoints (Just valid in replay mode, will not affect Linux kernel execution).

(gdb) b 375
Breakpoint 2 at 0xffffffff81179b75: file /build/buildd/linux-
3.2.0/fs/read_write.c, line 375.
(gdb) c
Continuing.

Breakpoint 2, vfs_read (file=0xffff8801f7bd4c00,
buf=0x7fff74e4edb0 <Address 0x7fff74e4edb0 out of bounds>,
count=16,
*   pos=0xffff8801f4b45f48) at /build/buildd/linux-*
3.2.0/fs/read_write.c:375

375   ret = rw_verify_area(READ, file, pos, count);
(gdb) s
rw_verify_area (read_write=0, file=0xffff8801f7bd4c00,
ppos=0xffff8801f4b45f48, count=16)
*   at /build/buildd/linux-3.2.0/fs/read_write.c:300*
300   inode = file->f_path.dentry->d_inode;

Use reverse-execution commands.

(gdb) rs

Breakpoint 2, vfs_read (file=0xffff8801f7bd4c00,
buf=0x7fff74e4edb0 <Address 0x7fff74e4edb0 out of bounds>,
count=16,
*   pos=0xffff8801f4b45f48) at /build/buildd/linux-*
3.2.0/fs/read_write.c:375
375   ret = rw_verify_area(READ, file, pos, count);
(gdb) rn
372   if (unlikely(!access_ok(VERIFY_WRITE, buf, count)))

GDB commands tstart, tfind or quit can auto close the replay mode.

Howto show a variable whose value has been optimized away

Sometimes, GDB will output some value like:

> *inode has been optimized out of existence.*
> *res has been optimized out of existence.*

That is because value of inode and res is optimized. Linux Kernel is built with -O2 so you will get this trouble sometimes.

There are 2 ways to handle it:

Update your GCC

The VTA branch http://gcc.gnu.org/wiki/Var_Tracking_Assignments was merged for GCC 4.5. This helps a lot with generating dwarf for previously "optimized out" values.

Get the way that access the variable that has been out through parse ASM code

Even if update the GCC to the newer version, you will still meet the issue. The main reason is the data is inside the registers but GCC doesn't put it to debug info. Then GDB just can output this variable has been optimized away.

But you can get where is the variable from ASM code and access it inside the tracepoint actions.

Following is a example that find variable "f" of function get_empty_filp and use it in tracepoint actions:

We want collect the value of "f" but looks it has been optimized away.

> *(gdb) list get_empty_filp*
> *...*
> *...*
> *...*
> *137   INIT_LIST_HEAD(&f->f_u.fu_list);*
> *138   atomic_long_set(&f->f_count, 1);*
> *139   rwlock_init(&f->f_owner.lock);*
> *140   spin_lock_init(&f->f_lock);*
> *141   eventpoll_init_file(f);*
> *(gdb)*
> *142   /* f->f_version: 0 */*
> *143   return f;*
> *(gdb) trace 143*
> *Tracepoint 1 at 0xffffffff8119b30e: file fs/file_table.c, line 143.*
> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*
> *End with a line saying just "end".*
> *>collect f*
> *`f' is optimized away and cannot be collected.*

Now use "disassemble /m" command get the ASM code and source line that have relation with "f" and parse them.

> *(gdb) disassemble /m get_empty_filp*
> *...*
> *...*
> *...*
> *125   f = kmem_cache_zalloc(filp_cachep, GFP_KERNEL);*
> *126   if (unlikely(!f))*
> *   0xffffffff8119b28c <+92>:   test   %rax,%rax*
> *   0xffffffff8119b292 <+98>:   je   0xffffffff8119b362*
> *<get_empty_filp+306>*
>
> *127   return ERR_PTR(-ENOMEM);*

```
    0xffffffff8119b362 <+306>:   mov    $0xfffffffffffffff4,%rax
    0xffffffff8119b369 <+313>:   jmp    0xffffffff8119b311
<get_empty_filp+225>
```

Code from "+98" to "+132" is not show in this part because they belong to other inline function. But you can get them with GDB command "disassemble get_empty_filp".

```
    0xffffffff8119b287 <+87>:    callq  0xffffffff81181cb0
<kmem_cache_alloc>
    0xffffffff8119b28c <+92>:    test   %rax,%rax
    0xffffffff8119b28f <+95>:    mov    %rax,%rbx
    0xffffffff8119b292 <+98>:    je     0xffffffff8119b362
<get_empty_filp+306>
    0xffffffff8119b298 <+104>:   mov    0xb4d406(%rip),%edx      #
0xffffffff81ce86a4 <percpu_counter_batch>
    0xffffffff8119b29e <+110>:   mov    $0x1,%esi
    0xffffffff8119b2a3 <+115>:   mov    $0xffffffff81c05340,%rdi
---Type <return> to continue, or q <return> to quit---
    0xffffffff8119b2aa <+122>:   callq  0xffffffff8130dd20
<__percpu_counter_add>
```

According to the ASM code you can see that return value of kmem_cache_alloc is inside $rax and its value is set to $rbx.

Looks $rbx has the value of "f". Let's check other ASM code.

```
128
129          percpu_counter_inc(&nr_files);
130          f->f_cred = get_cred(cred);
    0xffffffff8119b2b4 <+132>:   mov    %r12,0x70(%rbx)
```

Set a value to element of f, the ASM code is set value of $r12 to a address that base address is $rbx. It also looks like $rbx is "f".

```
131          error = security_file_alloc(f);
    0xffffffff8119b2b8 <+136>:   mov    %rbx,%rdi
    0xffffffff8119b2bb <+139>:   callq  0xffffffff8128ee30
<security_file_alloc>
```

```
132          if (unlikely(error)) {
    0xffffffff8119b2c0 <+144>:   test   %eax,%eax
    0xffffffff8119b2c2 <+146>:   jne    0xffffffff8119b36b
<get_empty_filp+315>
---Type <return> to continue, or q <return> to quit---
```

```
133              file_free(f);
134              return ERR_PTR(error);
```

```
0xffffffff8119b393 <+355>:  movslq -0x14(%rbp),%rax
0xffffffff8119b397 <+359>:  jmpq   0xffffffff8119b311
<get_empty_filp+225>

135        }
136
137        INIT_LIST_HEAD(&f->f_u.fu_list);
138        atomic_long_set(&f->f_count, 1);
139        rwlock_init(&f->f_owner.lock);
0xffffffff8119b2e4 <+180>:  movl   $0x100000,0x50(%rbx)

140        spin_lock_init(&f->f_lock);
0xffffffff8119b2c8 <+152>:  xor    %eax,%eax
0xffffffff8119b2d1 <+161>:  mov    %ax,0x30(%rbx)

141        eventpoll_init_file(f);
142        /* f->f_version: 0 */
143        return f;
0xffffffff8119b30e <+222>:  mov    %rbx,%rax
```

And after check other ASM code. You can make sure that $rbx is "f".

Then you can access "f" through access $rbx in tracepoint actions, for example:

```
(gdb) trace 143
Tracepoint 1 at 0xffffffff8119b30e: file fs/file_table.c, line 143.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
#collect f
>collect $rbx
#collect *f
>collect *((struct file *)$rbx)
#collect f->f_op
>collect ((struct file *)$rbx)->f_op
>end
```

How to get the function pointer point to

If the debug info of the function pointer is not optimized out

You can collect it directly and print what it point to. For example:

```
377                     count = ret;
378                     if (file->f_op->read)
379                             ret = file->f_op->read(file, buf, count, pos);
(gdb)
(gdb) trace 379
Tracepoint 1 at 0xffffffff81173ba5: file
/home/teawater/kernel/linux/fs/read_write.c, line 379.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
>collect file->f_op->read
>end
(gdb) tstart
(gdb) tstop
(gdb) tfind
(gdb) p file->f_op->read
$5 = (ssize_t (*)(struct file *, char *, size_t, loff_t *))
0xffffffff81173190 <do_sync_read>
#Then you know file->f_op->read point to do_sync_read.
```

If the debug info of the function pointer is optimized out

You can use tracepoint step to handle it. For example:

#Find out which instrunction that it is called.
(gdb) disassemble /rm vfs_read
379   ret = file->f_op->read(file, buf, count, pos);
  0xffffffff81173ba5 <+181>: 48 89 da   mov   %rbx,%rdx
  0xffffffff81173ba8 <+184>: 4c 89 e9   mov   %r13,%rcx
  0xffffffff81173bab <+187>: 4c 89 e6   mov   %r12,%rsi
  0xffffffff81173bae <+190>: 4c 89 f7   mov   %r14,%rdi
  *0xffffffff81173bb1 <+193>: ff d0   callq *%rax*
  0xffffffff81173bb3 <+195>: 48 89 c3   mov   %rax,%rbx
*(gdb) trace *0xffffffff81173bb1*
Tracepoint 1 at 0xffffffff81173bb1: file
/home/teawater/kernel/linux/fs/read_write.c, line 379.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
>while-stepping 1
 >collect $reg
 >end
>end
(gdb) tstart
(gdb) tstop
(gdb) tfind
#0 tty_read (file=0xffff88006ca74900, buf=0xb6b7dc <Address
0xb6b7dc out of bounds>, count=8176,
    ppos=0xffff88006e197f48) at
/home/teawater/kernel/linux/drivers/tty/tty_io.c:960
960   {
#Then you know file->f_op->read point to tty_read.

Please **note** that while-stepping will make tracepoint cannot use kprobes-optimization.

/sys/kernel/debug/gtpframe and offline debug

/sys/kernel/debug/gtpframe supplies trace frame in tfile format (GDB can parse it) when KGTP is stop.

In the PC that can run the GDB:

Change the "target remote XXXX" to

(gdb) target remote | perl ./getgtprsp.pl

After that, set tracepoint and start it as usual:

(gdb) trace vfs_readdir
Tracepoint 1 at 0xffffffff8114f3c0: file /home/teawater/kernel/linux-2.6/fs/readdir.c, line 24.
(gdb) actions
Enter actions for tracepoint 1, one per line.
End with a line saying just "end".
#If your GDB support tracepoint "printf" (see "Howto use tracepoint printf"), use it to show the value directly is better.
>collect $reg
>end
(gdb) tstart
(gdb) stop
(gdb) quit

Then you can find files gtpstart and gtpstop in current directory. Copy it to the machine that you want to debug.

In the debugged machine, copy the program "putgtprsp" and "gtp.ko" in the KGTP directory to this machine first. After insmod the gtp.ko:

Start the tracepoint:

./putgtprsp ./gtpstart

Stop the tracepoint:

./putgtprsp ./gtpstop

You can let Linux Kernel show the value directly, please see <u>Howto let tracepoint output value directly</u>.

If you want to save the value to the trace frame and parse later, you can use file "/sys/kernel/debug/gtpframe" that has the trace frame. Copy it to the PC

that has GDB.

Please **note** that some "cp" cannot handle it very well, please use "cat /sys/kernel/debug/gtpframe > ./gtpframe" to copy it.

In the PC that can run the GDB:

> *(gdb) target tfile ./gtpframe*
> *Tracepoint 1 at 0xffffffff8114f3dc: file /home/teawater/kernel/linux-2.6/fs/readdir.c, line 24.*
> *Created tracepoint 1 for target's tracepoint 1 at 0xffffffff8114f3c0.*
> *(gdb) tfind*
> *Found trace frame 0, tracepoint 1*
> *#0 vfs_readdir (file=0xffff880036e8f300, filler=0xffffffff8114f240 <filldir>, buf=0xffff880001e5bf38)*
> *   at /home/teawater/kernel/linux-2.6/fs/readdir.c:24*
> *24   {*

Please **note** that if you want connect KGTP from GDB in remote machine after use offline debug, you need "rmmod gtp" and "insmod gtp.ko" before call "nc".

How to use /sys/kernel/debug/gtpframe_pipe

This interface supplies same format trace frame with "gtpframe". But it can work when KGTP is running. After data is read, it will auto deleted from trace frame like "trace_pipe" of ftrace.

Get the frame info with GDB

#connect to the interface
(gdb) target tfile /sys/kernel/debug/gtpframe_pipe
#Get one trace frame entry
(gdb) tfind 0
Found trace frame 0, tracepoint 1
#Get the next one
(gdb) tfind
Target failed to find requested trace frame.
(gdb) tfind 0
Found trace frame 0, tracepoint 1

This way is better to work with python to parse Kernel. add-ons/hotcode.py is an example of python script.

Get the frame info with cat

sudo cat /sys/kernel/debug/gtpframe_pipe > g
Then all the trace frame will be saved in file "g".

Get the frame info with getframe

KGTP package include a program "getframe" can help you save the trace frame to files.

Following part is the help of it:

getframe -h
Get the trace frame of KGTP and save them in current
directory with tfile format.
Usage: ./getframe [option]

-g n   Set the minimum free size limit to n G.
When free size of current disk is smaller than n G,
./getframe will exit (-q) or wait some seconds (-w).
The default value of it is 2 G.

-q   Quit when current disk is smaller than
minimum free size limit (-g).

-w n   Wait n seconds when current disk is smaller
than minimum free size limit (-g).

-e n   Set the entry number of each tfile to n.
The default value of it is 1000.

-h   Display this information.

Use $pipe_trace

For the lock safe, KGTP will ignore the task that read the /sys/kernel/debug/gtpframe_pipe in default.

If you really need trace this task, and be sure that is safe. You can use following command before call "tstart":

    (gdb) tvariable $pipe_trace=1

Then KGTP will not ignore the task that read /sys/kernel/debug/gtpframe_pipe.

Use KGTP with user applications

KGTP can access the memory and trace user applications without stop it.

Let GDB connect KGTP for user applications

1) Open GDB without load any user applications.

2) If user applications is running on the current machine, use GDB command "target extended-remote /sys/kernel/debug/gtp" connect to KGTP. If user applications is running on the remote machine, use netcat like GDB on remote machine but replace "target remote" to "target extended-remote".

3) Load user applications (it must be built with GCC option "-g" to make it has debug information) with GDB command "file".

4) Use GDB command "attach pid" to attach the task's pid.

Then let GDB connect KGTP for user applications should be:

> *sudo gdb-release*
> *(gdb) target extended-remote /sys/kernel/debug/gtp*
> *Remote debugging using /sys/kernel/debug/gtp*
> *0x00000000 in ?? ()*
> *(gdb) file a.out*
> *A program is being debugged already.*
> *Are you sure you want to change the file? (y or n) y*
> *Reading symbols from /home/teawater/kernel/kgtp/a.out...done.*
> *(gdb) attach 15412*
> *A program is being debugged already. Kill it? (y or n) y*
> *Attaching to program: /home/teawater/kernel/kgtp/a.out, Remote target*
> *# Some version of GDB will output internal-error, please answer "n" to ignore it.*

Read memory of user applications directly

After GDB attach the user applications success, you can access the memory of this task with GDB commands "p" and "x". You can get help of these commands with GDB commands "help p" and "help x". For example:

```
(gdb) p c
$19 = 4460
(gdb) p &c
$21 = (int *) 0x601048 <c>
(gdb) x 0x601048
0x601048 <c>:   0x00001181
```

Trace user applications

KGTP use **uprobes** function of Linux kernel trace user applications, just Linux kernel 3.9 and later version support this function.

The build config of most Linux distributions's Linux kernel (3.9 and later) has opened **uprobes**.

For the Linux kernel that built by yourself:

> *Kernel hacking --->*
> *[*] Tracers --->*
> *[*]   Enable uprobes-based dynamic events*

If current Linux kernel **uprobes** is opened, you can set tracepoint according to GDB tracepoint after GDB attach the user applications success. For example:

> *(gdb) trace 14*
> *Tracepoint 1 at 0x400662: file /home/teawater/kernel/kgtp-misc/test.c, line 14.*
> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*
> *End with a line saying just "end".*
> *>collect $bt*
> *>collect c*
> *>end*
> *(gdb) tstart*
> *(gdb) tstatus*
> *Trace is running on the target.*
> *Collected 5 trace frames.*
> *Trace buffer has 20824428 bytes of 20828160 bytes free (0% full).*
> *Trace will stop if GDB disconnects.*
> *Not looking at any trace frame.*
> *(gdb) tstop*
> *(gdb) tfind*
> *Found trace frame 0, tracepoint 1*
> *#0 main (argc=1, argv=0x7fff5e878368, envp=0x7fff5e878378)*
> *at /home/teawater/kernel/kgtp-misc/test.c:14*
> *14   c += 1;*
> *(gdb) bt*
> *#0 main (argc=1, argv=0x7fff5e878368, envp=0x7fff5e878378)*
> *at /home/teawater/kernel/kgtp-misc/test.c:14*
> *(gdb) p c*
> *$7 = 36*

Please **note** that even if you just attch one of these tasks, user applications's tracepoint will be triggered by all tasks of a user application. (I think this is a very interesting feature of **uprobe**, so I didn't limit it in KGTP tracepoint.)

You can add $current_task_pid check to conditions of tracepoint to make tracepoint just be triggered by one of this task. Following example is set a tracepoint that just for task 985:

> *(gdb) trace 14*
> *Tracepoint 1 at 0x400662: file /home/teawater/kernel/kgtp-misc/test.c, line 14.*
> *(gdb) condition $bpnum ($current_task_pid == 985)*

And you can "collect $current_task_pid" in tracepoint actions to make sure which task triggers the tracepint. For example:

> *(gdb) trace 14*
> *Tracepoint 2 at 0x400662: file /home/teawater/kernel/kgtp-misc/test.c, line 14.*
> *(gdb) actions*
> *Enter actions for tracepoint 2, one per line.*
> *End with a line saying just "end".*
> *>collect $current_task_pid*
> *>collect c*
> *>end*
> *(gdb) tstart*
> *(gdb) tstatus*
> *Trace is running on the target.*
> *Collected 6 trace frames.*
> *Trace buffer has 20827776 bytes of 20828160 bytes free (0% full).*
> *Trace will stop if GDB disconnects.*
> *Not looking at any trace frame.*
> *(gdb) tstop*
> *(gdb) tfind*
> *Found trace frame 0, tracepoint 2*
> *#0 main (argc=<unavailable>, argv=<unavailable>, envp=<unavailable>) at /home/teawater/kernel/kgtp-misc/test.c:14*
> *14   c += 1;*
> *(gdb) p $current_task_pid*
> *$2 = 9983*
> *(gdb) tfind*
> *Found trace frame 1, tracepoint 2*
> *14   c += 1;*
> *(gdb) p $current_task_pid*
> *$3 = 9982*
> *(gdb)*

collect stack (for backtrace) of system from Linux kernel to user applications in tracepoint

$current is a special trace state variable that if the action of an tracepoint access it, this tracepint will access the values of the registers and the memory of current task instead of Linux kernel.

In general, the tracepoint will get the registers value of current task from **task_pt_regs**. Then collect **$current** in tracepoint actions will let this tracepoint access values of current task. For example:

> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*
> *End with a line saying just "end".*
> *>collect $current*
> *>collect $bt*
> *>end*

In addition, for some special function that its arguments include the pointer to the registers(for example: do_IRQ function of X86), tracepoint need get the registers from the arguments of fuction. Then set the pointer to $current will let this tracepoint get it. For example:

> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*
> *End with a line saying just "end".*
> *>teval $current=(uint64_t)regs*
> *>collect $bt*
> *>end*

$current_task_user is a special trace state variable that it is value will be true when current task is in user mode.

With these two trace state variables, you can use KGTP collect the stack(backtrace) of current task.

Following example show how we do backtrace(stack dump) from user space to Linux kernel:

> *#Connect to KGTP(same with prev section)*
> *(gdb) target extended-remote /sys/kernel/debug/gtp*
> *#Setup an tracepoint that collect the user space stack of task 18776.*
> *(gdb) trace vfs_read*
> *Tracepoint 1 at 0xffffffff8117a3d0: file*
> */home/teawater/kernel/linux/fs/read_write.c, line 365.*
> *(gdb) condition 1 ($current_task_user && $current_task_pid == 18776)*
> *(gdb) actions*
> *Enter actions for tracepoint 1, one per line.*

End with a line saying just "end".
>collect $current
>collect $bt
>end
#Setup a tracepoint that collect kernel space stack of task 18776.
(gdb) trace vfs_read
Note: breakpoint 1 also set at pc 0xffffffff8117a3d0.
Tracepoint 2 at 0xffffffff8117a3d0: file
/home/teawater/kernel/linux/fs/read_write.c, line 365.
(gdb) condition 2 ($current_task_user && $current_task_pid ==
18776)
(gdb) actions
Enter actions for tracepoint 2, one per line.
End with a line saying just "end".
>collect $bt
>end
(gdb) tstart
(gdb) tstop
#Following part is same with prev section, add a new inferior to
parse info of the user space program.
(gdb) add-inferior
Added inferior 2
(gdb) inferior 2
[Switching to inferior 2 [<null>] (<noexec>)]
(gdb) file gdb
Reading symbols from /usr/local/bin/gdb...done.
(gdb) attach 18776
#tracepoint 1 collect the user space stack.
(gdb) tfind
Found trace frame 0, tracepoint 1
#0 0x00007f77331d7d0f in __read_nocancel () from /lib/x86_64-
linux-gnu/libpthread.so.0
#This is the user space backtrace of task 18776.
(gdb) bt
#0 0x00007f77331d7d0f in __read_nocancel () from /lib/x86_64-
linux-gnu/libpthread.so.0
#1 0x000000000078e145 in rl_callback_read_char () at
../../src/readline/callback.c:201
#2 0x000000000069de79 in rl_callback_read_char_wrapper
(client_data=<optimized out>) at ../../src/gdb/event-top.c:169
#3 0x000000000069ccf8 in process_event () at ../../src/gdb/event-
loop.c:401
#4 process_event () at ../../src/gdb/event-loop.c:351
#5 0x000000000069d448 in gdb_do_one_event () at
../../src/gdb/event-loop.c:465
#6 0x000000000069d5d5 in start_event_loop () at
../../src/gdb/event-loop.c:490
#7 0x0000000000697083 in captured_command_loop

(data=<optimized out>) at ../../src/gdb/main.c:226
#8 0x0000000000695d8b in catch_errors (func=0x697070
<captured_command_loop>, func_args=0x0, errstring=0x14df99e
"",
   mask=6) at ../../src/gdb/exceptions.c:546
#9 0x00000000006979e6 in captured_main (data=<optimized
out>) at ../../src/gdb/main.c:1001
#10 0x0000000000695d8b in catch_errors (func=0x697360
<captured_main>,
   func@entry=<error reading variable: PC not available>,
func_args=0x7fff08afd5b0,
   func_args@entry=<error reading variable: PC not available>,
errstring=<unavailable>,
   errstring@entry=<error reading variable: PC not available>,
mask=<unavailable>,
   mask@entry=<error reading variable: PC not available>) at
../../src/gdb/exceptions.c:546
#11 <unavailable> in ?? ()
Backtrace stopped: not enough registers or memory available to
unwind further
#The tracepoint 2 collect the kernel space stack. So swith to
inferior 1 that load the kernel debug info.
(gdb) tfind
Found trace frame 1, tracepoint 2
#0 0xffffffff8117a3d0 in ?? ()
(gdb) inferior 1
[Switching to inferior 1 [Remote target]
(/home/teawater/kernel/b/vmlinux)]
[Switching to thread 1 (Remote target)]
#0 vfs_read (file=0xffff88021a559500, buf=0x7fff08afd31f
<Address 0x7fff08afd31f out of bounds>, count=1,
   pos=0xffff8800c47e1f48) at
/home/teawater/kernel/linux/fs/read_write.c:365
365   {
#This is the backtrace of kernel stack.
(gdb) bt
#0 vfs_read (file=0xffff88021a559500, buf=0x7fff08afd31f
<Address 0x7fff08afd31f out of bounds>, count=1,
   pos=0xffff8800c47e1f48) at
/home/teawater/kernel/linux/fs/read_write.c:365
#1 0xffffffff8117a59a in sys_read (fd=<optimized out>,
buf=0x7fff08afd31f <Address 0x7fff08afd31f out of bounds>,
   count=1) at /home/teawater/kernel/linux/fs/read_write.c:469
#2 <signal handler called>
#3 0x00007f77331d7d10 in ?? ()
#4 0x0000000000000000 in ?? ()

How to use add-ons/hotcode.py

This script can show the hottest code line in the Linux kernel or user space program through parse and record the pc address in the interrupt handler.

Please goto http://code.google.com/p/kgtp/wiki/hotcode see howto use it.

How to add plugin in C

KGTP support plugin that write in C. The plugin will be built as LKM

API

#include "gtp.h"
This header file include the API that plugin need.

*extern int gtp_plugin_mod_register(struct module *mod);*
*extern int gtp_plugin_mod_unregister(struct module *mod);*
These two functions register and unregister the plugin module. Then when
KGTP will add module usage count when it access the resource of the plugin
module.

*extern struct gtp_var *gtp_plugin_var_add(char *name, int64_t val,*
*struct gtp_var_hooks *hooks);*
This function add special trace state variable to the KGTP.

- **name** is the name of special trace state variable.

- **val** is initialization value of special trace state variable.

- **hooks** is the function pointers. The function pointers can be set to
  NULL if this function doesn't support.

- **Return** the gtp_var pointer if success. Get error will return error code
  that IS_ERR and PTR_ERR can handle.

struct gtp_var_hooks {
*   int   (*gdb_set_val)(struct gtp_trace_s *unused, struct gtp_var*
var,
*   int64_t val);*
*   int   (*gdb_get_val)(struct gtp_trace_s *unused, struct gtp_var*
var,
*   int64_t *val);*
*   int   (*agent_set_val)(struct gtp_trace_s *gts, struct gtp_var*
var,
*   int64_t val);*
*   int   (*agent_get_val)(struct gtp_trace_s *gts, struct gtp_var*
var,
*   int64_t *val);*
};
- **gdb_set_val** will be called when GDB set the value of TSV. Please note
  that TSV just can be set by GDB command "tvariable $xxx=1" and the

value just be sent to KGTP when GDB command "tstart".

- ◦ **unused** is unused. Just to make this pointer can share function with agent_set_val.
- ◦ **var** is the pointer that point to the gtp_var pointer. Then function of plugin can use it to figure out which TSV is accessed when TSVs share the function.
- ◦ **val** is the value that GDB set.
- ◦ **Return** return -1 if error. return 0 if success.

- • **gdb_get_val** will be called when GDB get the value of TSV. Please note that TSV get is different with TSV set. It can be gotten from KGTP anytime. And get its value just like get the value of GDB internal value. For example: "p $xxx".
  - ◦ **unused** is same with gdb_set_val.
  - ◦ **var** is same with gdb_set_val.
  - ◦ **val** is the pointer that use to return value.
  - ◦ **Return** is same with gdb_set_val.

- • **agent_set_val** will be called when tracepoint action(teval expr1, expr2, ...) set the value of TSV.
  - ◦ **gts** is pointer to the tracepoint session struct.
  - ◦ **var** is same with gdb_set_val.
  - ◦ **val** is the value the action set.
  - ◦ **Return** is same with gdb_set_val.

- • **agent_get_val** will be called when tracepoint action(collect expr1, expr2, ... or teval expr1, expr2, ...) get the value of TSV.
  - ◦ **gts** is same with agent_set_val.
  - ◦ **var** is same with gdb_set_val.
  - ◦ **val** is same with gdb_get_val.
  - ◦ **Return** is same with gdb_set_val.

*extern int gtp_plugin_var_del(struct gtp_var *var);*
When rmmod the plugin module, use this function remove the TSV that gtp_plugin_var_add add.

Example

plugin_example.c that in the KGTP directory is the example for KGTP plugin. You can use "make P=1" build it. It add 4 TSV to KGTP.

- **$test1** support nothing.

- **$test2** support be get and set by GDB or tracepoint action.

- **$test3** just support tracepoint action set. When set a value to it, it will look up a kernel symbol of this value and print it. For example "teval $test3=(int64_t)$rip".

- **$test4** just support tracepoint action set. When set a value to it, it will look up a kernel symbol of current tracepoint address and print it.

How to use

- insmod KGTP module according to <u>Insmod the KGTP module</u>.

- insmod plugin_example.ko

- Use GDB connect to KGTP and use it.

- Disconnect GDB. If option in <u>Do not stop tracepoint when the GDB disconnects</u> set to on, set it to off.

- rmmod plugin_example.ko

Please **note** that KGTP support add more than one plugin.

How to use performance counters

Performance counters are special hardware registers available on most modern CPUs. These registers count the number of certain types of hw events: such as instructions executed, cachemisses suffered, or branches mispredicted - without slowing down the kernel or applications. These registers can also trigger interrupts when a threshold number of events have passed - and can thus be used to profile the code that runs on that CPU.

The Linux Performance Counter subsystem called perf event can get the value of performance counter. You can access it through KGTP perf event trace state variables.

Please goto read the file tools/perf/design.txt in Linux Kernel to get more info about perf event.

Define a perf event trace state variable

Access an performance counter need define following trace state variable:

"pe_cpu_"+tv_name   Define the the CPU id of the performance counter.
"pe_type_"+tv_name   Define the the type of the performance counter.
"pe_config_"+tv_name   Define the the config of the performance counter.
"pe_en_"+tv_name   This the switch to enable or disable the performance counter.
The performance counter is disable in default.
"pe_val_"+tv_name   Access this variable can get the value of the performance counter.

Define a per_cpu perf event trace state variable

Define a per_cpu perf event trace state variable is same with define [Per_cpu trace state variables](#).

> *"p_pe_"+perf_event type+string+CPU_id*

Please **note** that if you define a per_cpu perf event trace state variable, you will not need define the cpu id("pe_cpu") because KGTP already get it.

The perf event type and config

The type of perf event can be:

0   PERF_TYPE_HARDWARE
1   PERF_TYPE_SOFTWARE
2   PERF_TYPE_TRACEPOINT
3   PERF_TYPE_HW_CACHE
4   PERF_TYPE_RAW
5   PERF_TYPE_BREAKPOINT

If the type is 0(PERF_TYPE_HARDWARE), the config can be:

0   PERF_COUNT_HW_CPU_CYCLES
1   PERF_COUNT_HW_INSTRUCTIONS
2   PERF_COUNT_HW_CACHE_REFERENCES
3   PERF_COUNT_HW_CACHE_MISSES
4   PERF_COUNT_HW_BRANCH_INSTRUCTIONS
5   PERF_COUNT_HW_BRANCH_MISSES
6   PERF_COUNT_HW_BUS_CYCLES
7   PERF_COUNT_HW_STALLED_CYCLES_FRONTEND
8   PERF_COUNT_HW_STALLED_CYCLES_BACKEND

If the type is 3(PERF_TYPE_HW_CACHE), the config need to divide to 3 parts:
First one is cache id, it need be << 0 before set to config:

0   PERF_COUNT_HW_CACHE_L1D
1   PERF_COUNT_HW_CACHE_L1I
2   PERF_COUNT_HW_CACHE_LL
3   PERF_COUNT_HW_CACHE_DTLB
4   PERF_COUNT_HW_CACHE_ITLB
5   PERF_COUNT_HW_CACHE_BPU

Second one is cache op id, it need be << 8 before set to config:

0   PERF_COUNT_HW_CACHE_OP_READ
1   PERF_COUNT_HW_CACHE_OP_WRITE
2   PERF_COUNT_HW_CACHE_OP_PREFETCH

Last one is cache op result id, it need be << 16 before set to config:

0   PERF_COUNT_HW_CACHE_RESULT_ACCESS
1   PERF_COUNT_HW_CACHE_RESULT_MISS

If you want get the perf count of PERF_COUNT_HW_CACHE_L1I(1), PERF_COUNT_HW_CACHE_OP_WRITE(1) and PERF_COUNT_HW_CACHE_RESULT_MISS(1), you can use:

(gdb) tvariable $pe_config_cache=1 | (1 << 8) | (1 << 16)

tools/perf/design.txt in Linux Kernel have more info about type and config of perf event.

Enable and disable all the perf event in a CPU with $p_pe_en

I think the best way that count a part of code with performance counters is enable all the count in the begin of the code and disable all of them in the end. You can do it with "pe_en". But if you have a lot of perf event trace state variables. That will make the tracepoint action very big. $p_pe_en is for this issue. You can enable all the perf event trace state variables in current CPU with following action:

>*teval $p_pe_en=1*

Disable them with set $p_pe_en to 0.

>*teval $p_pe_en=0*

GDB scripts to help with set and get the perf event trace state variables

Following is a GDB script define two commands dpe and spe to help define and show the perf event trace state variables.

You can put it to the ~/.gdbinit or your tracepoint script. Then you can use this two commands in GDB directly.

```
define dpe
  if ($argc < 2)
    printf "Usage: dpe pe_type pe_config [enable]\n"
  end
  if ($argc >= 2)
    eval "tvariable $p_pe_val_%d%d_c",$arg0, $arg1
    eval "tvariable $p_pe_en_%d%d_c",$arg0, $arg1
    set $tmp=0
    while $tmp<$cpu_number
      eval "tvariable $p_pe_type_%d%d_c%d=%d",$arg0, $arg1, $tmp, $arg0
      eval "tvariable $p_pe_config_%d%d_c%d=%d",$arg0, $arg1, $tmp, $arg1
      eval "tvariable $p_pe_val_%d%d_c%d=0",$arg0, $arg1, $tmp
      if ($argc >= 3)
        eval "tvariable $p_pe_en_%d%d_c%d=%d",$arg0, $arg1, $tmp, $arg2
      end
      set $tmp=$tmp+1
    end
  end
end

document dpe
Usage: dpe pe_type pe_config [enable]
end

define spe
  if ($argc != 2 && $argc != 3)
    printf "Usage: spe pe_type pe_config [cpu_id]\n"
  end
  if ($argc == 2)
    set $tmp=0
    while $tmp<$cpu_number
      eval "printf \"$p_pe_val_%%d%%d_c%%d=%%ld\\n\",$arg0, $arg1, $tmp, $p_pe_val_%d%d_c%d", $arg0, $arg1, $tmp
      set $tmp=$tmp+1
    end
```

```
  end
  if ($argc == 3)
    eval "printf \"$p_pe_val_%%d%%d_c%%d=%%ld\\n\",$arg0, $arg1,
$tmp, $p_pe_val_%d%d_c%d", $arg0, $arg1, $arg2
  end
end

document spe
Usage: spe pe_type pe_config [cpu_id]
end
```

Following is an example to use it get the performance counters of function tcp_v4_rcv:

```
#Connect to KGTP
(gdb) target remote /sys/kernel/debug/gtp
#Define 3 pe tvs for PERF_COUNT_HW_CPU_CYCLES,
PERF_COUNT_HW_CACHE_MISSES and
PERF_COUNT_HW_BRANCH_MISSES.
(gdb) dpe 0 0
(gdb) dpe 0 3
(gdb) dpe 0 5
#enable the performance counters of this CPU in the begin of this
function.
(gdb) trace tcp_v4_rcv
(gdb) action
>teval $p_pe_en=1
>end
#$kret make this hanler the end of function tcp_v4_rcv.
(gdb) trace *(tcp_v4_rcv)
(gdb) action
>teval $kret=0
#disable all performance counters of this CPU
>teval $p_pe_en=0
#Access the per cpu perf event tv will access to the current cpu pe
tv.
>collect $p_pe_val_00_0
>collect $p_pe_val_03_0
>collect $p_pe_val_05_0
#Set all the pe tv to 0
>teval $p_pe_val_00_0=0
>teval $p_pe_val_03_0=0
>teval $p_pe_val_05_0=0
>end
tstart
#Wait some time that current pc receive some tcp package.
(gdb) tstop
(gdb) tfind
```

(gdb) spe 0 0 $cpu_id
$p_pe_val_00_2=12676
(gdb) spe 0 3 $cpu_id
$p_pe_val_03_2=7
(gdb) spe 0 5 $cpu_id
$p_pe_val_05_2=97